

Adarsh Doli

adarshdoli257@gmail.com | 716-708-8999 | LinkedIn

Summary

Software Engineer with 2+ years of experience in full-stack development, cloud-native deployments, and AI/ML model support. Proficient in Java, Python, and SQL, with hands-on expertise in scalable application design and data pipeline automation. Strong understanding of Unix/Linux environments, computer vision, recommendations, accessibility, and reinforcement learning concepts.

Skills

- **Programming:** Java | Python | Go (Golang)
- **Frameworks & Libraries:** Flask | Spring | React | Node.js | TensorFlow | PyTorch
- **Databases:** PostgreSQL | MongoDB | NoSQL (DynamoDB) | Hadoop | Kafka
- **Cloud & DevOps:** AWS | Cloudflare | Kubernetes | Docker | Git | n8n | RESTful APIs | Linux
- **AI/ML:** Computer Vision | LangChain | Distributed Computing | RAG

Experience

Software Developer , Community Dreams Foundation - Remote

July 2025 – Present

- Designed a production RAG platform using Python, LangChain agents with MCP tools, ChromaDB/Qdrant retrieval, AWS S3/EKS/Elastic Beanstalk, and CloudWatch, launched with 99.99% uptime and 60% lower inference latency
- Developed LLM microservices in Java Spring Boot (REST), orchestrating retrieval, tool-use, inference, deployed on EC2+ALB, lifting pipeline throughput 45% and cutting p95 response time 30%
- Tuned Amazon RDS/Aurora (MySQL/PostgreSQL/Oracle) for RAG workloads, 6 schema/indexing for geospatial & time-series, achieving sub-second queries, 30% storage cost reduction, and reliability at 50k+ daily requests

Software Developer , Alorasoft Inc - Dallas, TX

June 2024 – Jan 2025

- Engineered backend services with Java and Spring Boot to power a cloud-based student learning platform, achieving 99.9% uptime and scalable performance.
- Integrated RESTful APIs across student dashboards, teacher portals, and assessments, delivering a 20% improvement in user workflow efficiency.
- Implemented real-time progress tracking with Kafka and WebSockets, boosting student engagement and raising course completion rates by 15%.
- Optimized PostgreSQL queries with indexes and schema redesign, cutting average response times by 40%
- Automated deployments with Docker and Kubernetes, reducing release cycles by 35% and enabling zero-downtime updates.

Associate Software Engineer, Hakri Software Services - Hyderabad, India

Jan 2023 – Jan 2024

- Developed a recommendation engine with Python, Pandas, and Scikit-learn, enhancing content relevance and driving a 12% increase in user retention
- Established version control for ML models and datasets with Git and DVC, ensuring 100% reproducibility and smooth deployment across environments.
- Designed an automated ETL pipeline with Apache Airflow and Docker to process nutritional data, cutting manual processing time by 70%.
- Deployed systems end-to-end on AWS (EC2, S3, RDS) and configured DevOps pipelines, strengthening CI/CD reliability.
- Built a customer portal with React.js and Flask APIs, improving load times by 30% through caching and CDN optimizations with CloudFront and Cloudflare.

Projects

Weather Streaming Project

- Designed Java microservices to process large-scale weather data in real time, using Apache Kafka and Apache Flink for event streaming.
- Boosted system performance with Redis and Ehcache caching, reducing average data access latency.
- Optimized PostgreSQL efficiency through indexing and stored procedures, improving data retrieval times by 50%.

Customer Support Chatbot

- Launched a customer service chatbot powered by LLaMA-3.1-8B, used LoRA fine-tuning and 4/8/16-bit quantization to cut response time by 40%.
- Implemented a scalable ML pipeline with Flask APIs and Ollama, containerized in Docker and deployed on AWS ECS Fargate, supporting up to 500 concurrent queries per day
- Automated CI/CD with AWS CodePipeline, reducing deployment errors by 35% and increasing iteration speed by 30%
- Instrumented performance with AWS CloudWatch, trimming inference latency by 22% and improving reliability via real-time alerts and logs

Education

University at Buffalo - SUNY Master of Science, Computer Science

Jan 2024 – June 2025

- GPA: 3.3/4.0